

CS 1810 Midterm Checklist

Introduction

Below, we outline some key concepts and ideas you should be comfortable with for the midterm. Under each topic, we have divided this into two or three sections:

1. Items to know
2. Things to be able to work through when given information/formulae
3. Items out of scope

Note that the third section is only defined for some topics. This list is not exhaustive, but we hope it is illustrative. We encourage you to review the course textbook, section materials, homework materials, and the practice exam problems for a full picture.

For emphasis: the midterm is not about memorization but will be designed to test your conceptual and analytical understanding.

1 Regression

1.1 Items to know

- What is the “bias trick” for rolling the bias into w ?
- What is the least squares loss function?
- How is w^* derived in a linear regression problem?
- What is a basis function, and why would we use it?
- What is regularization, and when do we use it? How does the loss function change for Lasso and ridge regression?
- What is the bias variance tradeoff, and how does it relate to overfitting, underfitting, regularization, and increasing the size of the training set?
- In Bayesian linear regression, how are priors related to regularization?
- What is a posterior distribution, posterior predictive distribution, and marginal likelihood?

1.2 Things to be able to work through when given information/equations

- You should be familiar with the mathematical techniques required to find w_{ridge} . Sketch the derivation, and how is this similar to the derivation for w^* for OLS?
- Given a likelihood function (you do *not* need to remember PDFs), how is the MLE computed?
- Given a Bayesian linear regression setup, how would the MAP be computed?

- Given expressions for conjugate distributions, and forms of posteriors, work through finding a posterior parameter or the posterior predictive for a given setting.

1.3 Items out of scope

- Second-order or other optimization methods (anything other than simple closed-form optimization and the gradient descent setup)
- Deriving conjugacies or using conjugacy statements from memory (e.g. Normal-Normal).

2 Classification

2.1 Items to know

- Role of hinge loss vs 0/1 loss vs logistic loss?
- What is the idea behind gradient descent (i.e. intuitively, why does it work)? What happens if the learning rate is too high? Too low?
- What are batch gradient descent and stochastic gradient descent?
- If the data are linearly separable, are there any guarantees about perceptron's behavior? What is the intuition behind this?
- What can we say about the shape of the decision boundary of a logistic regression?
- Can you explain all the parts of the likelihood function for two-class and multi-class logistic regression?
- How do we work with generative models for classification? What is the Naive Bayes model?
- What is the difference between a discriminative model like Logistic Regression and a generative model like Naive Bayes?

2.2 Things to be able to work through when given information/equations

- Gradients for negated log likelihood of logistic regression
- For a generative model and give class priors and class-conditional distributions, work with the (full) log-likelihood (using Lagrangian method where needed)
- Given parametric as well as non-parametric classifiers (such as kNN) and a particular dataset, how do we expect different classifiers to perform?
- Performing multiple, numerical steps of an iterative optimization such as gradient descent.
- Deriving basis functions to represent data in a space where it is separable (this is in scope for simple cases where such a transformation is easily discernible).

3 Neural Networks & Model Selection

3.1 Items to know

- Why do we need activation functions?
- What is backpropagation, and why is it useful?
- How can model selection methods such as cross-validation, and regularization be useful in supervised learning?
- How can neural networks be used for regression tasks and how can they be used for classification tasks?

3.2 Things to be able to work through when given information/equations

- Given a particular task for a neural network, what is a good choice for a loss function (e.g., least squares vs softmax)?
- Given a particular neural network architecture, how might changes to architecture (e.g. adding a layer, changing the structure) intuitively affect the model's performance?
- Given a simple neural network, determine whether the model is able to perform a particular classification task successfully.
- Given a simple (not necessarily named) loss function for a neural network, deduce why it is intuitive for the task at hand.
- Deriving backpropagation computations by hand

4 Support Vector Machines (SVMs)

4.1 Items to know

- What is the maximum margin classifier, and how does it relate to SVMs?
- What is the role of the hinge loss function in SVMs?
- What is the difference between hard-margin and soft-margin SVMs?
- How does the choice of the regularization parameter C affect the classifier?
- What is the kernel trick, and why is it useful?
- What are some common kernels (e.g. linear, polynomial, RBF), and how do they work?
- How does an SVM decision boundary compare to that of logistic regression?
- What are support vectors, and why are they important in SVMs?

4.2 Things to be able to work through when given information/equations

- Given a set of data points, determine the maximum margin classifier.
- Given a soft-margin SVM formulation, explain the role of slack variables and how they impact classification.
- Given a kernel function, compute the transformed feature space dimensions and explain its effect.
- Understand how SVMs handle non-linearly separable data using kernel functions.

4.3 Items out of scope

- Work through the dual formulation of an SVM and interpret the meaning of Lagrange multipliers.
- Solving the full quadratic optimization problem manually.
- Derivation of the Karush-Kuhn-Tucker (KKT) conditions for SVMs.